

A Tutorial on Programming Features in ATS

Hongwei Xi

hwxi AT cs DOT bu DOT edu

A Tutorial on Programming Features in ATS:
by Hongwei Xi

Copyright © 2010-2017 Hongwei Xi

All rights are reserved. Permission is granted to print this document for personal use.

Table of Contents

Preface	v
I. Basic Tutorial Topics	1
1. Syntax-Coloring for ATS code	1
2. Filename Extensions	3
3. File Inclusion inside ATS Code	5
4. Fixity Declarations	7
5. The Program Entry Point: mainats	9
6. Tail-Recursive Call Optimization	11
7. Mutual Tail-Recursion	13
8. Metrics for Termination Verification.....	15
Primitive Recursion.....	15
General Recursion	15
Mutual Recursion	16
Termination Checking at Run-time.....	16
9. Higher-Order Functions.....	19
10. Parametric Polymorphism	21
11. Printf-like Functions	23
12. Functional Lists.....	25
13. Persistent Arrays	27
14. Persistent References	29
15. Call-by-Reference	31
16. Lazy Evaluation.....	33
II. Advanced Tutorial Topics.....	35
17. Cast Functions	35
18. Stack Allocation at Run-Time	37

Preface

This tutorial covers a variety of issues that a programmer typically encounters when programming in ATS. It is written mostly in the style of learn-by-examples. Although it is possible to learn programming in ATS by studying the tutorial (if the reader is familiar with ML and C), I consider the book *Introduction to Programming in ATS*¹ a far more appropriate resource for someone to acquire a view of ATS in a coherent and systematic manner. Of course, there will also be considerable amount of overlapping between these two books. The primary purpose of the tutorial is to bring more insights into a rich set of programming features in ATS and also demonstrate through concrete examples that these features can be made of effective use in the construction of high-quality programs.

Notes

1. <http://www.ats-lang.org/DOCUMENT/INTPROGINATS/HTML/book1.html>

Chapter 1. Syntax-Coloring for ATS code

The syntax of ATS is highly involved, which can be a daunting obstacle for beginners trying to read and write ATS code. In order to alleviate this problem, I may employ colors to differentiate various syntactical entities in ATS code. The convention I adopt for coloring ATS syntax is given as follows:

- The keywords in ATS are all colored black (and possibly written in bold face).
- The comments in ATS are all colored gray.
- The code in the statics of ATS is colored blue.
- The code in the dynamics of ATS is colored red unless it represents proofs, for which the color dark green is used.
- The external code (in C) is colored deep brown.

Please find an example of ATS code on-line¹ that involves all of the syntax-coloring mentioned above.

Notes

1. http://www.ats-lang.org/DOCUMENT/TUTORIALATS/CODE/fact_dats.html

Chapter 2. Filename Extensions

In ATS, the filename extensions *.sats* and *.dats* are reserved to indicate static and dynamic files, respectively. As these two extensions have some special meaning attached to them, which can be interpreted by the command **atscc**, they should not be replaced arbitrarily.

A static file may contain sort definitions, datasort declarations, static definitions, abstract type declarations, exception declarations, datatype declarations, macro definition, interfaces for dynamic values and functions, etc. In terms of functionality, a static file in ATS is somewhat similar to a header file (with the extension *.h*) in C or an interface file (with the extension *.mli*) in OCaml.

A dynamic file may contain everything in a static file. In addition, it may also contain definitions for dynamic values and functions. However, interfaces for functions and values in a dynamic file should follow the keyword **extern**, which on the other hand should not be present when such interfaces are declared in a static file. For instance, the following syntax declares interfaces (or types) in a static file for a value named **pi** and a function named **area_of_circle**:

```
val pi : double
fun area_of_circle (radius: double): double
```

When the same interfaces are declared in a *dynamic* file, the following slightly different syntax should be used:

```
extern val pi : double
extern fun area_of_circle (radius: double): double
```

As a convention, we often use the filename extension *.cats* to indicate that a file contains some C code that is supposed to be combined with ATS code in certain manner. There is also the filename extension *.hats*, which we occasionally use for a file that should be included in ATS code stored in another file. However, the use of these two filename extensions are not mandatory, that is, they can be replaced if needed or wanted.

Chapter 3. File Inclusion inside ATS Code

As is in C, file inclusion inside ATS code is done through the use of the directive *#include*. For instance, the following line indicates that a file named *foobar* is included, that is, this line is to be replaced with the content of the file *foobar*:

```
#include "foobar.hats"
```

Note that the included file is parsed according to the syntax for statics or dynamics depending on whether the file is included in a static or dynamic file. As a convention, the name of an included file often ends with the extension *.hats*.

A common use of file inclusion is to keep some constants, flags or parameters being defined consistently across a set of files. For instance, the file `prelude/params.hats`¹ serves such a purpose. File inclusion can also be used to emulate (in a limited but rather useful manner) functors supported in languages such as SML and OCaml.

Notes

1. <http://www.ats-lang.org/DOCUMENT/ANAIIRIATS/prelude/params.hats>

Chapter 4. Fixity Declarations

Given a function f , the standard syntax for applying f to an argument v is $f(v)$; for two arguments $v1$ and $v2$, the syntax is $f(v1, v2)$. However, it is allowed in ATS to use infix notation for a binary function application, and prefix/postfix notation for a unary function application.

Each identifier in ATS can be assigned one of the following fixities: *prefix*, *infix* and *postfix*. The fixity declarations for many commonly used identifiers can be found in `prelude/fixity.ats`¹. Often, the name *operator* is used to refer to an identifier that is assigned a fixity. For instance, the following syntax declares that `+` and `-` are infix operators of a precedence value equal to 50:

```
infixl 50 + -
```

After this declaration, we can write an expression like `1 + 2 - 3`, which is parsed into `-(+(1, 2), 3)` in terms of the standard syntax for function application.

The keyword `infixl` indicates that the declared infix operators are left-associative. For right-associative and non-associative infix operators, please use the keywords `infixr` and `infix`, respectively. If the precedence value is omitted in a fixity declaration, it is assumed to be equal to 0.

We can also use the following syntax to declare that `iadd`, `fadd`, `padd` and `uadd` are left-associative infix operators with a precedence value equal to that of the operator `+`:

```
infixl (+) iadd fadd padd uadd
```

This is useful as it is difficult in practice to remember the precedence values of (a large collection of) declared operators. Sometimes, we may need to specify that the precedence value of one operator in relation to that of another one. For instance, the following syntax declares that `opr2` is a left-associative infix operator and its precedence value is that of `opr1` plus 10:

```
infixl (opr1 + 10) opr2
```

If the plus sign (+) is changed to the minus sign (-), then the precedence value of `opr2` is that of `opr1` minus 10.

We can also turn an identifier `opr` into a non-associative infix operator (of precedence value 0) by putting the backslash symbol (`\`) in front of it. For instance, the expression `exp1 \opr exp2` stands for `opr (exp1, exp2)`, where `exp1` and `exp2` refer to some expressions, either static or dynamic.

The syntax for declaring (unary) prefix and postfix operators are similar. For instance, the following syntax declares that `~` and `?` are prefix and postfix operators of precedence values 61 and 69, respectively:

```
prefix 61 ~
postfix 69 ?
```

As an example, a postfix operator is involved in the following 3-line program:

```
postfix (imul + 10) !!
extern fun !! (x: int): int
implement !! (x) = if x >= 2 then x * (x - 2)!! else 1
```

For a given occurrence of an operator, we can deprive it of its assigned fixity status by simply putting the keyword `op` in front of it. For instance `1 + 2 - 3` can be written as `op-(op+ (1, 2), 3)`. It is also possible to (permanently) deprive an operator of its assigned fixity status. For instance, the following syntax does so to the operators `iadd`, `fadd`, `padd` and `uadd`:

```
nonfix iadd fadd padd uadd
```

Lastly, please note that each fixity declaration is only effective within its own legal scope.

Notes

1. <http://www.ats-lang.org/DOCUMENT/ANAIIRIATS/prelude/fixity.ats>

Chapter 5. The Program Entry Point: mainats

There are two special functions of the name `main_void` and `main_argc_argv` that are given the following interfaces:

```
fun main_void (): void = "mainats"
overload main with main_void

fun main_argc_argv {n:int | n >= 1}
  (argc: int n, argv: &(@[string][n])): void = "mainats"
overload main with main_argc_argv
```

The symbol `main` is overloaded with both of these functions. In addition, the global name `mainats` is assigned to both of them. When a function in ATS is translated into one in C, the global name of the function, if ever assigned, is used to refer to its translation in C.

The interface for `main_argc_argv` indicates that the function takes as its arguments an integer `argc` greater than 0 and an array `argv` of size `argc` in which each element is a string, and returns no value. The syntax `argv: &(@[string][n])` means that `argv` is a call-by-reference argument. If we followed the like of syntax in C++, then this would be written as something like `&argv: @[string][n]`.

To turn ATS source code into an executable, the function is required to be present in the C code translated from the ATS code (as it is called within the `main` function in C). Normally, this means that either `main_void` or `main_argc_argv` needs to be implemented in the ATS code (that is to be turned into an executable). However, in certain situations, it may make sense to implement `mainats` in C directly. Note that the interface for `mainats` in C is:

```
extern ats_void_type mainats (ats_int_type, ats_ptr_type) ;
```

where `ats_void_type`, `ats_int_type` and `ats_ptr_type` are just aliases for `void`, `int` and `void*`, respectively.

As an example, the following ATS program echos onto the standard output the given command line:

```
implement
main (argc, argv) = let
  fun loop {n,i:nat | i <= n} // [loop] is tail-recursive
    (i: int i, argc: int n, argv: &(@[string][n])): void =
      if i < argc then begin
        if i > 0 then print (' '); print argv.[i]; loop (i+1, argc, argv)
      end // end of [if]
  // end of [loop]
in
  loop (0, argc, argv); print_newline ()
end // end of [main]
```

If `mainats` needs to be implemented in C, the proof function `main_dummy` should be implemented as follows:

```
implement main_dummy () = ()
```

This implementation is solely for telling the ATS compiler that `mainats` is expected to be implemented in C directly so that the compiler can generate proper code to handle the situation. As an example, I present as follows a typical scenario in GTK+ programming, where the function `gtk_init` needs to be called to modify the arguments passed from the given command line:

Chapter 5. The Program Entry Point: *mainats*

```
//
// some function implemented in ATS
//
extern fun main_work // implemented elsewhere
  {n:pos} (argc: int n, argv: &([string][n])): void = "main_work"
// end of [main_work]

%{^
// %{^ : the embedded C code is placed at the top
extern ats_void_type mainats (ats_int_type, ats_ptr_type) ;
%} // end of [%{^]

implement main_dummy () = () // indicating [mainats] being implemented in C

%{$
// %{$ : the embedded C code is placed at the bottom
ats_void_type
mainats (
  ats_int_type argc, ats_ptr_type argv
) {
  gtk_init ((int*)&argc, (char ***)&argv) ; main_work (argc, argv) ; return ;
} /* end of [mainats] */

%} // end of [%{$]
```


Chapter 6. Tail-Recursive Call Optimization

Tail-recursion is of crucial importance in functional programming as loops in imperative programming are often implemented as tail-recursive functions.

Suppose that a function *foo* calls another function *bar*, that is, there is a call to *bar* appearing in the body of *foo*, where *foo* and *bar* may actually be the same function. If the return value of the call to *bar* also happens to be the return value of *foo*, then this call is often referred to as a *tail-call*. If *foo* and *bar* are the same, then this call is a (recursive) self tail-call. For instance, there are two recursive calls in the body of the following defined function **f91**:

```
fun f91 (n: int): int = if n > 100 then n - 10 else f91 (f91 (n+11))
```

where the outer call to **f91** is a tail-call while the inner one is not. If each self call in the body of a function is a tail-call, then this function is tail-recursive.

It is arguably that the single most important optimization performed by the ATS compiler is the translation of every self tail-call into a direct (local) jump. This optimization effectively turns every tail-recursive function into the equivalent of a loop, guaranteeing that a fixed amount of stack space is adequate for executing each call to the function.

Let us now see some examples that can help further illustrate the notion of tail-recursive call optimization. The following recursive function **sum1** sums up all the natural numbers less than or equal to a given integer parameter **n**:

```
fun sum1 (n: int): int = if n > 0 then n + sum1 (n-1) else 0
```

Clearly, **sum1** is not tail-recursive as the only self call in its body is not a tail-call. The counterpart of **sum1** in C can essentially be given as follows:

```
int sum1 (int n) {  
    return (n > 0) ? n + sum1 (n-1) : 1 ;  
} // end of [sum1]
```

When applied to an integer **n**, the following defined function **sum2** also sums up all the natural numbers less than or equal to **n**:

```
fn sum2 (n: int): int = let  
    fun loop (n: int, res: int): int =  
        if n > 0 then loop (n-1, res + n) else res  
    // end of [loop]  
in  
    loop (n, 0)  
end // end of [sum2]
```

The inner function **loop** in the definition of **sum2** is tail-recursive. The stack space consumed by **loop** is constant, that is, it is independent of the values of the arguments of **loop**. Essentially, the ATS compiler translates the definition of **sum2** into the following C code:

```
int sum2_loop (int n, int res) {  
    loop:  
    if (n > 0) {  
        res = res + n ; n = n - 1; goto loop;  
    } else {  
        // do nothing  
    }  
    return res ;  
}
```

```
}  
  
int sum2 (int n) { return sum2_loop (n, 0) ; }
```

Translating **sum1** into **sum2** is a fairly straightforward process. However, it can be highly involved, sometimes, to turn a non-tail-recursive implementation into an equivalent one that is tail-recursive.

Chapter 7. Mutual Tail-Recursion

Mutually tail-recursive functions are commonly encountered in practice. Assume that *foo* and *bar* are two mutually defined functions. In the body of either *foo* or *bar*, a tail-call to *foo* or *bar* is referred to as a mutually tail-recursive call. If every call to *foo* or *bar* in the bodies of *foo* and *bar* are tail-call, then *foo* and *bar* are mutually tail-recursive. Mutual recursion involving more functions can be defined similarly. As an example, the following two functions *isEvn* and *isOdd* are defined mutually recursively:

```
fun isEvn (n: int): bool = if n > 0 then isOdd (n-1) else true
and isOdd (n: int): bool = if n > 0 then isEvn (n-1) else false
```

The call to *isOdd* in the body of *isEvn* is a mutually tail-recursive call, and the call to *isEvn* in the body of *isOdd* is also a mutually tail-recursive call. Therefore, *isEvn* and *isOdd* are mutually tail-recursive.

In order to turn mutually recursive tail-calls into jumps, the bodies of the involved mutually recursive functions need to be combined. The keyword *fn** in ATS is introduced precisely for this purpose. For instance, let us replace *fun* with *fn** in the code above:

```
fn* isEvn (n: int): bool = if n > 0 then isOdd (n-1) else true
and isOdd (n: int): bool = if n > 0 then isEvn (n-1) else false
```

Then the following C code is essentially what is generated from compiling these two mutually defined functions:

```
bool isEvnisOdd (int tag, int n) {
    bool res ;

    switch (tag) {
        0: goto isEvn ;
        1: goto isOdd ;
        default : exit (1) ;
    }

    isEvn: if (n > 0) { n = n - 1; goto isodd; } else { res = true; goto done; }
    isOdd: if (n > 0) { n = n - 1; goto isEvn; } else { res = false; goto done; }

    done: return res ;

} /* end of [isEvnisOdd] */

bool isEvn (int n) { return isEvnisOdd (0, n) ; }
bool isOdd (int n) { return isEvnisOdd (1, n) ; }
```

Note that mutually recursive functions can be combined in such a manner only if they all have the same return type. In the above case, both *isEvn* and *isOdd* have the same return type *bool*.

When translating C code involving embedded loops, we often encounter the need for mutual tail-recursion. For example, the following C code prints out some ordered pairs of digits:

```
int main (
    int argc, char *argv[]
) {
    int i, j ;
    for (i = 0; i <= 9; i += 1) {
        for (j = i; j <= 9; j += 1) {
```

```

        if (i < j) printf (" ") ; printf ("%i, %i", i, j) ;
    } /* for */
    printf ("
") ;
} /* for */
return 0 ;
}

```

A straightforward translation of the C code into ATS (in the style of functional programming) can be done as follows:

```

implement
main (argc, argv) = let
  fn* loop1
    {i:nat | i <= 10} (i: int i): void =
    if i <= 9 then loop2 (i, i) else ()
  // end of [loop1]

  and loop2
    {i,j:nat | i <= 9; j <= 10} (i: int i, j: int j): void =
    if j <= 9 then (
      if i < j then (
        print " "; printf ("%i, %i", @i, j); loop2 (i, j+1)
      ) // end of [if]
    ) else (
      print_newline (); loop1 (i+1)
    ) // end of [if]
  // end of [loop2]
in
  loop1 (0)
end // end of [main]

```

Evidently, **loop1** and **loop2** are defined mutually tail-recursively. The use of the keyword **fn*** ensures that these two functions are translated into the equivalent of two loops in C, which require only a fixed amount of stack space to run.

Chapter 8. Metrics for Termination Verification

ATS provides a simple means for the programmer to verify the termination of recursively defined functions by supplying proper termination metrics. This is an indispensable feature for supporting the paradigm of programming with theorem-proving as proof functions, namely, functions representing proofs, must be proven to be terminating.

A termination metric is a tuple (M_1, \dots, M_n) of natural numbers, where $n \geq 0$. We use the standard well-founded lexicographical ordering on natural numbers to order such tuples.

Primitive Recursion

The kind of recursion in the following implementation of the factorial function is primitive recursion:

```
fun fact {n:nat} .<n>.
  (n: int n): int = if n > 0 then n * fact (n-1) else 1
// end of [fact]
```

The special syntax `.<n>` indicates that the metric supplied for verifying the termination of the defined function is a singleton tuple (n) . In the definition of `fact`, the metric for the recursive call to `fact` is $(n-1)$, which is strictly less than the original metric (n) . Therefore, the defined function `fact` is terminating.

General Recursion

We implement as follows a function `gcd` that computes the greatest common divisor of two given positive integers:

```
//
// computing the greatest common divisors of two positive ints
//
fun gcd
  {m,n:int | m > 0; n > 0} .<m+n>.
  (m: int m, n: int n): [r:nat | 1 <= r; r <= min(m, n)] int r =
    if m > n then gcd (m - n, n)
    else if m < n then gcd (m, n - m)
    else m
// end of [gcd]
```

The syntax `.<m+n>` indicates that the termination metric $(m+n)$ is supplied to verify that the defined function `gcd` is terminating. In the definition of `gcd`, the termination metric for the first recursive call to `gcd` is $(m-n)+n=m$, which is strictly less than the original termination metric $m+n$ (as n is positive); the termination metric for the second recursive call to `gcd` is $m+(n-m)=n$, which is also strictly less than the original termination metric $m+n$ (as m is positive). Thus, `gcd` is a terminating function.

As another example, we implement as follows the Ackermann's function, which is famous for being recursive but not primitive recursive:

```
//
// [ack] implements the Ackermann's function
//
fun ack {m,n:nat} .<m, n>.
  (m: int m, n: int n): Nat =
    if m > 0 then
      if n > 0 then ack (m-1, ack (m, n-1)) else ack (m-1, 1)
    else n+1 // end of [if]
// end of [ack]
```

The syntax $\langle m, n \rangle$ indicates that the termination metric is a pair of natural numbers: (m, n) . Note that the standard lexicographical ordering on natural numbers is employed to compare such metrics. To verify that `ack` is terminating, we need to solve the following constraints:

- $(m-1, k)$ is less than (m, n) under the assumption $m > 0$, where k can be any natural number.
- $(m, n-1)$ is less than (m, n) under the assumption $m > 0$ and $n > 0$.
- $(m-1, 1)$ is less than (m, n) under the assumption $m > 0$.

As all of these constraints can be readily solved, we conclude that `ack` is a terminating function.

Mutual Recursion

When mutually recursive functions are to be verified, the termination metrics for these functions, which are tuples of natural numbers, must be of the same tuple length. We give a simple example as follows:

```
fun isEvn
  {n:nat} .<2*n+2>. (n: int n): bool =
    if n > 0 then ~(isOdd n) else true // end of [if]
// end of [isEvn]

and isOdd
  {n:nat} .<2*n+1>. (n: int n): bool =
    if n > 0 then isEvn (n-1) else false // end of [if]
// end of [isOdd]
```

Clearly, we can also verify the termination of these two functions by using the metrics $\langle n, 1 \rangle$ and $\langle n, 0 \rangle$ for `isEvn` and `isOdd`, respectively.

Termination Checking at Run-time

Suppose that `foo` and `bar` are declared as follows:

```
fun foo ():<> void and bar ():<> void
```

Moreover, suppose that the following implementation of `foo` is given in a file named `foo.dats`:

```
implement foo () = $Bar.bar ()
```

while the following implementation of `bar` is given in another file named `bar.dats` that is different from `foo.dats`:

```
implement bar () = $Foo.foo ()
```

Clearly, neither `foo` nor `bar` is terminating. In practice, it is difficult to resolve this issue of calling cycles among functions by solely relying on termination metrics. Instead, `atscc` can generate run-time code for detecting calling cycles among functions if the flag `-D_ATS_PROOFCHECK` is present. For instance, if `foo.dats` and `bar.dats` are compiled as follows:

```
atscc -D_ATS_PROOFCHECK foo.dats and bar.dats
```

then a run-time error is to be reported to indicate a calling cycle when either **foo.dats** or **bar.dats** is loaded dynamically.

Chapter 9. Higher-Order Functions

A higher-order function is one that takes another function as its argument. Let us use BT to range over base types such as **char**, **double**, **int** and **string**. A simple type T is formed according to the following inductive definition:

- BT is a simple type.
- $(T_1, \dots, T_n) \rightarrow T_0$ is a simple type if T_0, T_1, \dots, T_n are simple types.

Let *order* be a function from simply types to natural numbers defined as follows:

- $order(BT) = 0$
- $order((T_1, \dots, T_n) \rightarrow T_0) = \max(order(T_0), 1 + order(T_1), \dots, 1 + order(T_n))$

Given a function *f* of some simple type *T*, we say that *f* is a *n*th-order function if $order(T) = n$. For instance, a function of the type $(int, int) \rightarrow int$ is 1st-order, and a function of the type $int \rightarrow (int \rightarrow int)$ is also 1st-order, and a function of the type $((int \rightarrow int), int) \rightarrow int$ is 2nd-order. In practice, most higher-order functions are 2nd-order.

As an example, we implement as follows a 2nd-order function **find_root** that takes as its only argument a function *f* from integers to integers and searches for a root of *f* by enumeration:

```
fn find_root
  (f: int -<cloref1> int): int = let
    fun aux (
      f: int -<cloref1> int, n: int
    ) : int =
      if f (n) = 0 then n else (
        if n <= 0 then aux (f, ~n + 1) else aux (f, ~n)
      ) // end of [if]
  in
    aux (f, 0)
  end // end of [find_root]
```

The function **find_root** computes the values of *f* at 0, 1, -1, 2, -2, etc. until it finds the first integer *n* in this sequence that satisfies $f(n) = 0$.

As another example, we implement as follows the famous Newton-Raphson's method for finding roots of functions on reals:

```
typedef
  fdouble = double -<cloref1> double
  //
  macdef epsilon = 1E-6 (* precision *)
  //
  // [f1] is the derivative of [f]
  //
  fn newton_raphson (
    f: fdouble, f1: fdouble, x0: double
  ) : double = let
    fun loop (
      f: fdouble, f1: fdouble, x0: double
    ) : double = let
      val y0 = f x0
    in
      if abs (y0 / x0) < epsilon then x0 else
        let val y1 = f1 x0 in loop (f, f1, x0 - y0 / y1) end
      // end of [if]
    end // end of [loop]
```

```
in
  loop (f, f1, x0)
end // end of [newton_raphson]
```

We can now readily implement square root function and the cubic root function based on **newton_raphson**:

```
// square root function
fn sqrt (c: double): double =
  newton_raphson (lam x => x * x - c, lam x => 2.0 * x, 1.0)
// cubic root function
fn cbrt (c: double): double =
  newton_raphson (lam x => x * x * x - c, lam x => 3.0 * x * x, 1.0)
```

Higher-order functions can be of great use in supporting a form of code sharing that is both common and flexible. As function arguments are often represented as heap-allocated closures that can only be reclaimed through garbage collection (GC), higher-order functions are used infrequently, if ever, in a setting where GC is not present. In ATS, linear closures, which can be manually freed if needed, are available to support higher-order functions in the absence of GC, making it possible to employ higher-order functions extensively in systems programming (where GC is unavailable or simply disallowed). The details on linear closures are to be given elsewhere.

Chapter 10. Parametric Polymorphism

Parametric polymorphism (or polymorphism for short) offers a flexible and effective approach to supporting code reuse. For instance, given a pair (v1, v2) where v1 and v2 are a boolean a character, respectively, the function `swap_bool_char` defined below returns a pair (v2, v1):

```
fun swap_bool_char (xy: (bool, char)): (char, bool) = (xy.1, xy.0)
```

Suppose that integer values need to be swapped as well. This leads to the implementation of the following function `swap_int_int`:

```
fun swap_int_int (xy: (int, int)): (int, int) = (xy.1, xy.0)
```

The code duplication between `swap_bool_char` and `swap_int_int` is obvious, and it can be easily avoided by implementing a function template as follows:

```
fun{a,b:t@type} swap (xy: (a, b)): (b, a) = (xy.1, xy.0)
```

Now the functions `swap_bool_char` and `swap_int_int` can simply be replaced with `swap<bool,char>` and `swap<int,int>`, respectively. The function template `swap` cannot be compiled into binary object code directly as the sizes of type variables `a` and `b` are unknown: The special sort `t@type` is for classifying types whose sizes are unspecified. If `swap<T1,T2>` is used for some types T1 and T2 of known sizes, then an instantiation of `swap` is created where type variables `a` and `b` are replaced with T1 and T2, respectively, and the instantiation is compiled into binary object code. For those know the feature of templates in C++, this should sound rather familiar.

In contrast to `swap`, `swap_type_type` is defined below as a polymorphic function (rather than a function template):

```
fun swap_type_type {a,b:type} (xy: (a, b)): (b, a) = (xy.1, xy.0)
```

This function can be compiled into binary object code as the sizes of type variables `a` and `b` are known: The special sort `type` is for classifying types of size equal to exactly one word, that is, the size of a pointer. For example, the size of a string is one word, and the size of any declared datatype is also one word. Given strings `str1` and `str2`, an application of `swap_type_type` to `str1` and `str2` can be written as follows:

```
swap_type_type {string,string} (str1, str2)
```

where the expression `{string,string}` is often referred to as a static argument. As in this case, most static arguments do not have to be supplied explicitly since they can be automatically inferred. However, such static arguments, if given, can often greatly enhance the quality and precision of the error messages reported in case of type-checking failure.

Chapter 11. Printf-like Functions

The `printf` function in C is variadic, that is, the arity of the function is indefinite. In ATS, there is a function of the same name that is essentially the counterpart of the `printf` function in C.

The following interface is assigned to `printf` in ATS:

```
fun printf {ts:types} (fmt: printf_c ts, arg: ts): void
```

We use `printf_c` for a type constructor that takes a list of types to form a type for format strings (in C). For instance, `printf_c(char, double, int)` is a type for format strings that require a character, a double, and an integer to be supplied. Given a character `c`, a double `d` and an integer `i`, `@(c, d, i)` is an argument of types `(char, double, int)`, and the following expression is well-typed in ATS:

```
printf ("c = %c and d = %f and i = %i", @(c, d, i))
```

The type of the format string `"c = %c and d = %f and i = %i"` is computed to be `printf_c(char, double, int)` and then `@(c, d, i)` is checked to be of the type `(char, double, int)`. Note that a format string must be a constant in order for its type to be computed during typechecking.

As an example, we present as follows a program that prints out a multiplication table for single digits:

```
#define N 9

implement
main () = let
  fun loop1
    {i:nat | i <= N}
    (i: int i): void =
      if i < N then loop2 (i+1, 0) else ()
  // end of [loop1]
  and loop2
    {i,j:nat | i <= N; j <= i}
    (i: int i, j: int j): void =
      if j < i then let
        val () = if (j > 0) then print ' '
        val () = printf ("%1d*%1d=%2.2d", @(j+1, i, (j+1) * i))
      in
        loop2 (i, j+1)
      end else let
        val () = print_newline () in loop1 (i)
      end // end of [if]
  // end of [loop2]
in
  loop1 (0)
end // end of [main]
```

This programs generates the following expected output:

```
1*1=01
1*2=02 2*2=04
1*3=03 2*3=06 3*3=09
1*4=04 2*4=08 3*4=12 4*4=16
1*5=05 2*5=10 3*5=15 4*5=20 5*5=25
1*6=06 2*6=12 3*6=18 4*6=24 5*6=30 6*6=36
1*7=07 2*7=14 3*7=21 4*7=28 5*7=35 6*7=42 7*7=49
1*8=08 2*8=16 3*8=24 4*8=32 5*8=40 6*8=48 7*8=56 8*8=64
1*9=09 2*9=18 3*9=27 4*9=36 5*9=45 6*9=54 7*9=63 8*9=72 9*9=81
```

Please find a few other functions declared in `prelude/SATS/printf.sats`¹ that are similar to `printf`.

Notes

1. <http://www.ats-lang.org/DOCUMENT/ANAIIRIATS/prelude/SATS/printf.sats>

Chapter 12. Functional Lists

Lists are by far the most commonly used data structure in functional programming. We say that a data structure is **functional** if it is heap-allocated and immutable and can only be freed through garbage collection (GC). In contrast, a data structure is said to be linear if it is either stack-allocated or heap-allocated and can be freed by the user as well as by the GC.

The datatype for functional lists in ATS is (essentially) declared as follows:

```
datatype list (a:t@ype, int) =
  | {n:nat} list_cons (a, n+1) of (a, list (a, n))
  | list_nil (a, 0) of ()
// end of [list]
```

There are two data constructors associated with **list**: **list_nil** forms an empty list and **list_cons** for a list of a given head and tail. Given a type T and an integer I, the type **list(T, I)** is for lists of length I in which each element is of type T. Note that the sort **t@ype** indicates that the element type of a list can be unboxed (i.e., flat).

Often the following abbreviations are introduced for the list constructors so as to make the code involving list-processing less verbose:

```
#define nil list_nil
#define cons list_cons
#define :: list_cons // [::] is an infix operator
```

As an example of list creation, the following expression evaluates to a list consisting of integers 1, 2 and 3:

```
cons (1, cons (2, cons (3, nil ()))) // [nil ()] can be replaced with [nil]
```

Clearly, this kind of syntax is a bit unwieldy if longer lists need to be constructed. The following alternatives can also be used to create lists:

```
val xs = '[1, 2, 3] // the first character is quote (')
val xs = $lst (1, 2, 3) // this is equivalent to '[1, 2, 3]
val xs = $lst {Nat} (1, 2, 3) // [Nat] is given as the type for the list elements
```

The interfaces for various functions on lists can be found in `prelude/SATS/list.sats`¹.

Let us now see some list-processing code in ATS. The following program implements a function template that computes the length of a given list:

```
fun{a:t@ype}
length {n:nat} .<n>.
  (xs: list (a, n)): int n =
    case+ xs of _ :: xs => 1 + length xs | nil () => 0
// end of [length]
```

As this is not a tail-recursive implementation, the function **length** may have difficulty handling long lists (e.g., those containing more than 1 million elements). A tail-recursive implementation of **length** that can handle lists of any length is given as follows:

```
fun{a:t@ype}
length {n:nat} .<>.
  (xs: list (a, n)): int n = let
    fun loop {i,j:nat} .<i>.
      (xs: list (a, i), j: int j): int (i+j) =
```

```
      case+ xs of _ :: xs => loop (xs, j+1) | nil () => j
    // end of [loop]
  in
    loop (xs, 0)
  end // end of [length]
```

Let us see another example. The following function **append** returns the concatenation of two given lists:

```
fun{a:t@type}
append {m,n:nat} .<m>. (
  xs: list (a, m), ys: list (a, n)
) : list (a, m+n) =
  case+ xs of
  | cons (x, xs) => cons (x, append (xs, ys)) | nil () => ys
// end of [append]
```

This is not a tail-recursive implementation. As a consequence, **append** may have difficulty handling a case where its first argument is of a large length (e.g., 1 million). Can **append** be given a tail-recursive implementation in ATS? The answer is affirmative. For instance, a tail-recursive implementation of **append** is available in `prelude/DATS/list.dats`². As the implementation makes use of linear types, it is to be explained elsewhere.

Notes

1. <http://www.ats-lang.org/DOCUMENT/ANAIRIATS/prelude/SATS/list.sats>
2. <http://www.ats-lang.org/DOCUMENT/ANAIRIATS/prelude/DATS/list.dats>

Chapter 13. Persistent Arrays

A persistent array of size n is just n heap-allocated cells (or references) in a row. It is persistent in the sense that the memory allocated for the array cannot be freed manually. Instead, it can only be reclaimed through garbage collection (GC).

Given a viewtype VT , the type for arrays containing n values of viewtype VT is `array(VT, n)`. Note that arrays in ATS are the same as those in C: There is no size information attached them. The interfaces for various functions on arrays can be found in `prelude/SATS/array.sats`¹.

There are various functions for array creation. For instance, the following two are commonly used:

```
fun(a:t@type)
array_make_elt
  {n:nat} (asz: size_t n, elt: a):<> array (a, n)
// end of [array_make_elt]

fun(a:t@type)
array_make_lst {n:nat}
  (asz: size_t n, xs: list (a, n)):<> array (a, n)
// end of [array_make_lst]
```

Applied to a size and an element, `array_make_elt` returns an array of the given size in which each cell is initialized with the given element. Applied to a size and a list of elements, `array_make_lst` returns an array of the given size in which each cell is initialized with the corresponding element in the given list.

For reading from and writing to an array, the function templates `array_get_elt` and `array_set_elt` can be used, respectively, which are assigned the following interfaces:

```
fun(a:t@type)
array_get_elt_at {n:int}
  {i:nat | i < n} (A: array (a, n), i: size_t i):<!ref> a

fun(a:t@type)
array_set_elt_at {n:int}
  {i:nat | i < n} (A: array (a, n), i: size_t i, x: a):<!ref> void
```

Given an array A , an index i and a value v , `array_get_elt_al(A, i)` and `array_set_elt_at(A, i, v)` can be written as $A[i]$ and $A[i] := v$, respectively.

As an example, the following function template reverses the content of a given array:

```
fun(a:t@type)
array_reverse {n:nat} (
  A: array (a, n), n: size_t n)
) : void = let
  fun loop {i: nat | i <= n} .<n-i>.
    (A: array (a, n), n: size_t n, i: size_t i): void =
      if i < n/2 then let
        val tmp = A[i]
      in
        A[i] := A[n-1-i]; A[n-1-i] := tmp; loop (A, n, i+1)
      end else () // end of [if]
    // end of [loop]
  in
    loop (A, n, 0)
end // end of [array_reverse]
```

If the test $i < n/2$ is changed to $i \leq n/2$, a type-error is to be reported. Why? The reason is that $A[n-1-i]$ becomes out-of-bounds array subscripting in the case where n and i both equal zero. Given that it is very unlikely to encounter a case where an array of size 0 is involved, a bug like this, if not detected early, can be buried so scarily deep!

The careful reader may have already noticed that the sort `t@type` is assigned to the template parameter `a`. In other words, the above implementation of `array_reverse` cannot handle a case where the values stored in a array are of a linear type. The reason for choosing the sort `t@type` is that both `array_get_elt_at` and `array_set_elt_at` can only be applied an array containing values of a nonlinear type. In the following implementation, the template parameter is given the sort `viewt@type` so that an array containing values of a linear type can be handled:

```
fun{a:viewt@type}
array_reverse {n:nat} (
  A: array (a, n), n: size_t (n)
) : void = let
  fun loop {i: nat | i <= n} .<n-i>.
    (A: array (a, n), n: size_t n, i: size_t i): void =
      if i < n/2 then let
        val () = array_exch (A, i, n-1-i) in loop (A, n, i+1)
      end else () // end of [if]
    // end of [loop]
  in
    loop (A, n, 0)
  end // end of [array_reverse]
```

The interface for the function template `array_exch` is given below:

```
fun{a:viewt@type}
array_exch {n:nat}
  (A: array (a, n), i: sizeLt n, j: sizeLt n):<!ref> void
// end of [array_exch]
```

Note that `array_exch` can not be implemented in terms of `array_get_elt_at` and `array_set_elt_at` (unless some form of type-unsafe coding is employed). The curious reader can find its type-safe implementation in `prelude/DATS/array.dats`², which is based on a corresponding operation for linear arrays.

Notes

1. <http://www.ats-lang.org/DOCUMENT/ANAIIRIATS/prelude/SATS/array.sats>
2. <http://www.ats-lang.org/DOCUMENT/ANAIIRIATS/prelude/DATS/array.dats>

Chapter 14. Persistent References

A reference is essentially a heap-allocated array of size 1. It is persistent in the sense that the memory allocated for storing the content of a reference cannot be freed manually. Instead, it can only be reclaimed through garbage collection (GC).

Given a viewtype VT, the type for references to values of viewtype VT is `ref(VT)`. For convenience, the type constructor `ref` is declared to be abstract in ATS. However, it can be defined as follows:

```
typedef ref (a:viewt@type) = [l:addr] (vbox (a @ 1) | ptr 1)
```

The interfaces for various functions on references can be found in `prelude/SATS/reference.sats`¹.

For creating a reference, the function template `ref_make_elt` of the following interface can be called:

```
fun(a:viewt@type) ref_make_elt (x: a):<> ref a
```

For reading from and writing to a reference, the function templates `ref_get_elt` and `ref_set_elt` can be used, respectively, which are assigned the following interfaces:

```
fun(a:t@type) ref_get_elt (r: ref a):<!ref> a
fun(a:t@type) ref_set_elt (r: ref a, x: a):<!ref> void
```

Note that the symbol `!ref` indicates that these functions incur the so-called `ref`-effect when evaluated. Given a reference `r` and a value `v`, `ref_get_elt(r)` and `ref_set_elt(r, v)` can be written as `!r` and `!r := v`, respectively.

A reference is typically employed to record some form of persistent state. For instance, following is such an example:

```
local
//
// [ref] is a shorthand for [ref_make_elt]
//
val count = ref<int> (0)

in // in of [local]

fun getNewName
  (prfx: string): string = let
    val n = !count
    val () = !count := n + 1
    val name = sprintf ("%s%i", @(prfx, n))
  in
    string_of_strptr (name)
  end // end of [getNewName]

end // end of [local]
```

The function `getNewName` is called to generate fresh names. As the integer content of the reference `count` is updated whenever a call to `getNewName` is made, each name returned by `getNewName` is guaranteed to have not generated before. Note that each string returned by `sprintf` is a linear one (of the type `strptr`) and the cast function `string_of_strptr` is called to turn it into a nonlinear one. There is no run-time cost associated with such a call as every call to a cast function is always a no-op at run-time.

References are commonly misused in practice. The following program is often written by a beginner of functional programming who has already learned (some) imperative programming:

```
fun fact
  (n: int): int = let
    val res = ref<int> (1)
    fun loop (n: int):<cloref1> void =
      if n > 0 then !res := n * !res else ()
    val () = loop (n)
  in
    !res
end // end of [fact]
```

The function `fact` is written in such a style as somewhat a direct translation of the following C code:

```
int fact (int n) {
  int res = 1 ;
  while (n > 0) res = n * res ;
  return res ;
}
```

In the ATS implementation of `fact`, `res` is a heap-allocated reference and it becomes garbage (waiting to be reclaimed by the GC) after a call to `fact` returns. On the other hand, the variable `res` in the C implementation of `fact` is stack-allocated (or it can even be mapped to a machine register), and there is no generated garbage after a call to `fact` returns. A proper translation of the C implementation in ATS can actually be given as follows, which makes no use of references:

```
fun fact
  (n: int): int = let
    fun loop (n: int, res: int): int =
      if n > 0 then loop (n, n * res) else res
    // end of [loop]
  in
    loop (n, 1)
  end // end of [fact]
```

Unless strong justification can be given, making extensive use of (dynamically created) references is often a sure sign of poor coding style.

Notes

1. <http://www.ats-lang.org/DOCUMENT/ANAIATS/prelude/SATS/reference.sats>

Chapter 15. Call-by-Reference

The feature of call-by-reference in ATS is similar to the corresponding one in C++. What is special in ATS is the way in which this feature is handled by the type system. In general, if f is given a type of the following form for some viewtypes $VT1$ and $VT2$:

```
(..., &VT1 >> VT2, ...) -> ...
```

then a function call $f(\dots, lval, \dots)$ on some left-value $lval$ of the viewtype $VT1$ is to change the viewtype of $lval$ into $VT2$ upon its return. In the case where $VT1$ and $VT2$ are the same, $\&VT1 \gg VT2$ may simply be written as $\&VT1$.

As an example, an implementation of the factorial function is given as follows that makes use of call-by-reference:

```
fun fact (x: int): int = let
  fun loop {l:addr} (x: int, res: &int): void =
    if x > 0 then (res := res * x; loop (x-1, res)) else ()
  var res: int = 1 // [res] is a variable!
in
  loop (x, res); res
end // end of [fact]
```

Note that if the line for introducing the variable res in the above implementation is replaced with the following one:

```
val res: int = 1 // [res] is no longer a variable but a value!
```

then a type error is to be reported as res is no longer a left-value when it is passed as an argument to `loop`.

In functional programming, optional types are often used for error-handling. For instance, the following function `divopt` returns a value of the type `Option(int)` that either contains the result of the division or indicates a case of division-by-zero:

```
fun divopt
  (x: int, y: int): Option (int) =
  if y != 0 then Some (x/y) else None ()
// end of [divopt]
```

Given that a value of the form `Some(v)` is heap-allocated, the memory for storing it can only be reclaimed by garbage collection (GC). In other words, the memory is leaked if GC is not available. To address the issue of error-handling in the absence of GC, we can employ call-by-reference as is shown below:

```
fun diverr (
  x: int, y: int, err: &int
) : int =
  if y != 0 then x/y else (err := err+1; 0(*meaningless*))
// end of [diverr]
```

We can tell whether division-by-zero has occurred by comparing the values of `err` before and after a call to `diverr`. This style of error-handling is commonly seen in code written in languages like C.

Chapter 16. Lazy Evaluation

Though ATS is a language based on eager call-by-value evaluation, it also allows the programmer to perform lazy call-by-need evaluation. In ATS, there is a special language construct `$delay` that can be used to delay or suspend the evaluation of an expression (by forming a thunk) and a special function `lazy_force` that can be called to resume a suspended evaluation (represented by a thunk).

There is a special type constructor `lazy` of the sort `(t@type) => type` in ATS, which forms a (boxed) type when applied to a type. On one hand, given an expression `exp` of type `T`, `$delay(exp)` forms a value of the type `lazy(T)` that represents the suspended evaluation of `exp`. On the other hand, given a value `v` of the type `lazy(T)` for some type `T`, `lazy_force(v)` resumes the suspended evaluation represented by `v` and returns a result of type `T`. The interface for the function template `lazy_force` is given as follows:

```
fun{a:t@type} lazy_force (x: lazy a):<!laz> a
```

Note that the symbol `!laz` indicates a form of effect associated with lazy-evaluation. For cleaner syntax, the special prefix operator `!` in ATS is overloaded with `lazy_force`.

In prelude/SATS/lazy.sats¹, the following datatype `stream_con` and `stream` are declared mutually recursively for representing (lazy) streams:

```
datatype stream_con (a:t@type+) =  
  | stream_nil (a) of () | stream_cons (a) of (a, stream a)  
where stream (a:t@type) = lazy (stream_con a)
```

Also, a number of common functions on streams are declared in prelude/SATS/lazy.sats² and implemented in prelude/DATS/lazy.dats³.

The following code gives a standard implementation of the sieve of Eratosthenes. We first construct a stream of all the integers starting from 2 that are ordered ascendingly; we keep the first element of the stream and remove all of its multiples; we repeat this process on the rest of the stream recursively. The final stream then consists of all the prime numbers ordered ascendingly.

```
#define nil stream_nil  
#define cons stream_cons  
#define :: stream_cons  
  
typedef N2 = [n:int | n >= 2] int n  
val N2s: stream N2 = from 2 where {  
  fun from (n: N2):<!laz> stream N2 = $delay (n :: from (n+1))  
}  
  
fun sieve  
  (ns: stream N2):<!laz> stream N2 = let  
    // [val-] means no warning message from the compiler  
    val- n :: ns = !ns  
  in  
    $delay (n :: sieve (stream_filter_cloref<N2> (ns, lam x => x nmod n > 0)))  
  end // end of [sieve]  
  
val primes: stream N2 = sieve N2s  
  
//  
// Finding the nth prime where counting starts from 0  
//  
fn nprime {n: nat} (n: int n): N2 = stream_nth (primes, n)
```

The function template `stream_filter_cloref` is of the following interface:

```
fun(a:t@type)
stream_filter_cloref
  (xs: stream a, p: a -<cloref,!laz> bool):<!laz> stream a
```

It is called to construct a stream consisting of all the elements in a given stream that satisfy a given predicate.

We give another example of lazy-evaluation as follows, which demonstrates an interesting approach to computing Fibonacci numbers:

```
val one = int64_of_int 1

val // the following values are defined mutually recursively
rec fibs_1: stream int64 = $delay (one :: fibs_2) // fib1, fib2, ...
and fibs_2: stream int64 = $delay (one :: fibs_3) // fib2, fib3, ...
and fibs_3: stream int64 = ( // fib3, fib4, ...
  stream_map2_fun<int64,int64><int64> (fibs_1, fibs_2, lam (x, y) => x + y)
)
// find the nth Fibonacci number
fn nfib {n:pos} (n: int n): int64 = stream_nth (fibs_1, n-1)
```

The function template `stream_map2_fun` is assigned the following interface:

```
fun(a1,a2:t@type){b:t@type}
stream_map2_fun
  (xs1: stream a1, xs2: stream a2, f: (a1, a2) -<!laz> b):<!laz> stream b
```

Given two streams `xs1` and `xs2` and a binary function `f`, `stream_map2_fun` forms a stream such that the n th element in it, if it exists, equals $f(x_1, x_2)$, where x_1 and x_2 are the n th elements in `xs1` and `xs2`, respectively.

Notes

1. <http://www.ats-lang.org/DOCUMENT/ANAIIRIATS/prelude/SATS/lazy.sats>
2. <http://www.ats-lang.org/DOCUMENT/ANAIIRIATS/prelude/SATS/lazy.sats>
3. <http://www.ats-lang.org/DOCUMENT/ANAIIRIATS/prelude/DATS/lazy.dats>

Chapter 17. Cast Functions

A cast function in ATS is equivalent to the identify function in terms of dynamic semantics. A call to such a function is evaluated at compile-time, and its argument is returned. For instance, we have the following commonly used cast functions:

```
castfn int1_of_int (x: int):<> [n:nat] int n
castfn string1_of_string (x: string):<> [n:nat] string n
```

Note that the keyword **castfn** is for introducing cast functions.

Let us now see a more interesting use of casting functions. The following declared function interface is intended for concatenating a list of lists:

```
extern fun{a:t@type} list_concat (xss: List (List a)): List a
```

Assume that we would like to verify that the concatenation of a list of lists yields a list whose length equals the sum of the lengths of the lists in the given list of lists. This, for instance, can be done as follows by introducing a datatype constructor **lstlst**.

```
datatype lstlst (a:t@type+, int, int) =
  | {m,t:nat} {n:nat}
    lstlst_cons (a, m+1, t+n) of (list (a, n), lstlst (a, m, t))
  | lstlst_nil (a, 0, 0) of ()
// end of [lstlst]

fun{a:t@type} _concat {m,t:nat} .<m>.
  (xss: lstlst (a, m, t)):<> list (a, t) = case+ xss of
  | lstlst_cons (xs, xss) => list_append (xs, _concat<a> xss)
  | lstlst_nil () => list_nil ()
// end of [_concat]
```

Given a type **T** and integers **I** and **J**, the type **lstlst(T, I, J)** is for a list of lists such that the length of the list is **I** and each element in the list is a list of values of the type **T** and the sum of the lengths of these elements equals **J**. The function **list_concat** is the same as the function **_concat** in terms of dynamic semantics, and it can be implemented as follows:

```
implement{a}
list_concat (xss) =
  _concat<a> (lstlst_of_listlist xss) where {
    castfn lstlst_of_listlist
      {m:nat} .<m>. (xss: list (List a, m))
      :<> [t:nat] lstlst (a, m, t) = case+ xss of
      | list_cons (xs, xss) => lstlst_cons (xs, lstlst_of_listlist xss)
      | list_nil () => lstlst_nil ()
  } // end of [list_concat]
```

Given **lstlst_of_listlist** being implemented as a casting function, the implementation of **list_concat** is equivalent to the following one in terms of dynamic semantics:

```
implement{a}
list_concat (xss) = _concat (xss) // this one does not typecheck
```


Chapter 18. Stack Allocation at Run-Time

In ATS, there is support for allocating memory at run-time in the stack frame of the calling function, and it is guaranteed by the type system of ATS that the memory thus allocated cannot be accessed once the calling function returns.

In the following contrived example, the implemented function `name_of_month_1` allocates in its stack frame an array of size 12 that is initialized with the names of 12 months, and then returns the name of the i th month, where i is an integer between 1 and 12:

```
fn name_of_month_1
  {i:int | 1 <= i; i <= 12} (i: int i): string = let
  var !p_arr with pf_arr = @[string](
    "Jan", "Feb", "Mar", "Apr", "May", "Jun", "Jul", "Aug", "Sep", "Oct", "Nov", "Dec"
  ) // end of [var]
in
  p_arr->[i-1] // it can also be written as !p_arr[i-1]
end // end of [name_of_month_1]
```

The following syntax means that the starting address of the allocated array is stored in `p_arr` while the view of the array is stored in `pf_arr`:

```
var !p_arr with pf_arr = @[string](
  "Jan", "Feb", "Mar", "Apr", "May", "Jun", "Jul", "Aug", "Sep", "Oct", "Nov", "Dec"
) // end of [var]
```

This allocated array is initialized with the strings representing the names of the 12 months: "Jan", "Feb", "Mar", "Apr", "May", "Jun", "Jul", "Aug", "Sep", "Oct", "Nov", "Dec".

A variant of the function `name_of_month_1` is implemented as follows:

```
fn name_of_month_2
  {i:int | 1 <= i; i < 12}
  (i: int i): string = let
  var !p_arr with pf_arr = @[string][12]("")
  val () = p_arr->[0] := "Jan"
  val () = p_arr->[1] := "Feb"
  val () = p_arr->[2] := "Mar"
  val () = p_arr->[3] := "Apr"
  val () = p_arr->[4] := "May"
  val () = p_arr->[5] := "Jun"
  val () = p_arr->[6] := "Jul"
  val () = p_arr->[7] := "Aug"
  val () = p_arr->[8] := "Sep"
  val () = p_arr->[9] := "Oct"
  val () = p_arr->[10] := "Nov"
  val () = p_arr->[11] := "Dec"
in
  p_arr->[i-1]
end // end of [name_of_month_2]
```

The following syntax means that the function `name_of_month_2` allocates a string array of size 12 in its stack frame and initializes the array with the empty string:

```
var !p_arr with pf_arr = @[string][12]("")
```

The starting address and the view of the allocated array are stored in `p_arr` and `pf_arr`, respectively. If the following syntax is used:

```
var !p_arr with pf_arr = @[string][12]()
```

then the allocated array is uninitialized, that is, the view of the proof `pf_arr` is `[string?][12] @ p_arr` (instead of `[string][12] @ p_arr`).

When higher-order functions are employed in systems programming, it is often desirable to form closures in the stack frame of the calling function so as to avoid the need for memory allocation on heap. In the following example, the implemented function `print_month_name` forms a closure in its stack frame, which is then passed to a higher-order function `foreach_array_ptr_clo`:

```
fn print_month_names () = let
  var !p_arr with pf_arr = @[string](
    "Jan", "Feb", "Mar", "Apr", "May", "Jun", "Jul", "Aug", "Sep", "Oct", "Nov", "Dec"
  ) // end of [var]
//
  var !p_clo with pf_clo = @lam // this closure is stack-allocated
    (i: sizeLt 12, x: &string): void =<clo1> (if i > 0 then print ", "; print x)
  // end of [var]
  val () = array_ptr_iforeach_clo<string> (!p_arr, !p_clo, 12)
//
  val () = print_newline ()
in
  // empty
end // end of [print_month_names]
```

Note that the keyword `@lam` (instead of `lam`) is used here to indicate that the closure is constructed in the stack frame of the function `print_month_names`.